

In []:

Dataset Overview

HAM10000 ("Human Against Machine with 10000 training images") dataset - a large collection of multi-source dermatoscopic images of pigmented lesions

The dermatoscopic images are collected from different populations, acquired and stored by different modalities. The final dataset consists of 10015 dermatoscopic images.

It has 7 different classes of skin cancer which are listed below :

- Melanocytic nevi
- Melanoma
- Benign keratosis-like lesions
- Basal cell carcinoma
- Actinic keratoses
- Vascular lesions
- Dermatofibroma

Importing libraries

In [1]:

```
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
from imblearn.over_sampling import RandomOverSampler
import numpy as np
from sklearn.model_selection import train_test_split
import os, cv2
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, Flatten, Dense, MaxPool2D, Activation, Batch
Normalization
from sklearn.metrics import classification_report, accuracy_score
```

Reading the Data

In [2]:

```
data = pd.read_csv('/kaggle/input/skin-cancer-mnist-ham10000/hmnist_28_28_RGB.csv')
data.head()
```

Out[2]:

	pixel0000	pixel0001	pixel0002	pixel0003	pixel0004	pixel0005	pixel0006	pixel0007	pixel0008	pixel0009	...	pixel2343	pix
0	192	153	193	195	155	192	197	154	185	202	...	173	
1	25	14	30	68	48	75	123	93	126	158	...	60	
2	192	138	153	200	145	163	201	142	160	206	...	167	
3	38	19	30	95	59	72	143	103	119	171	...	44	
4	158	113	139	194	144	174	215	162	191	225	...	209	

5 rows × 2353 columns



Data Preprocessing

Data Cleaning

In [3]:

```
data['label'].unique()
```

Out[3]:

```
array([2, 4, 3, 6, 5, 1, 0])
```

In [4]:

```
y = data['label']  
x = data.drop(columns = ['label'])
```

In [5]:

```
data.isnull().sum().sum() #no null values present
```

Out[5]:

```
0
```

In [6]:

```
meta_data = pd.read_csv('/kaggle/input/skin-cancer-mnist-ham10000/HAM10000_metadata.csv')  
meta_data.head()
```

Out[6]:

	lesion_id	image_id	dx	dx_type	age	sex	localization
0	HAM_0000118	ISIC_0027419	bkl	histo	80.0	male	scalp
1	HAM_0000118	ISIC_0025030	bkl	histo	80.0	male	scalp
2	HAM_0002730	ISIC_0026769	bkl	histo	80.0	male	scalp
3	HAM_0002730	ISIC_0025661	bkl	histo	80.0	male	scalp
4	HAM_0001466	ISIC_0031633	bkl	histo	75.0	male	ear

In [7]:

```
meta_data['dx'].unique()
```

Out[7]:

```
array(['bkl', 'nv', 'df', 'mel', 'vasc', 'bcc', 'akiec'], dtype=object)
```

In [8]:

```
y = data['label']  
x = data.drop(columns = ['label'])
```

In [9]:

```
data.isnull().sum().sum() #no null values present
```

Out[9]:

```
0
```

In [10]:

```
meta_data = pd.read_csv('/kaggle/input/skin-cancer-mnist-ham10000/HAM10000_metadata.csv')  
meta_data.head()
```

Out[10]:

	lesion_id	image_id	dx	dx_type	age	sex	localization
0	HAM_0000118	ISIC_0027419	bkl	histo	80.0	male	scalp
1	HAM_0000118	ISIC_0025030	bkl	histo	80.0	male	scalp
2	HAM_0002730	ISIC_0026769	bkl	histo	80.0	male	scalp
3	HAM_0002730	ISIC_0025661	bkl	histo	80.0	male	scalp
4	HAM_0001466	ISIC_0031633	bkl	histo	75.0	male	ear

In [11]:

```
meta_data['dx'].unique()
```

Out[11]:

array(['bkl', 'nv', 'df', 'mel', 'vasc', 'bcc', 'akiec'], dtype=object)


Exploratory Data Analysis

In [12]:

```
sns.countplot(x = 'dx', data = meta_data)
plt.xlabel('Disease(Classes)', size=12)
plt.ylabel('Frequency', size=12)
plt.title('Frequency Distribution of Classes')
```

Out[12]:

Text(0.5, 1.0, 'Frequency Distribution of Classes')




In [13]:

```
sns.set_style('whitegrid')
colors = ['#87ace8', '#e377c2', 'green']
fig, axes = plt.subplots(figsize=(8,8))


ax = sns.countplot(x='sex', data=meta_data, palette = 'Paired')
for container in ax.containers:
    ax.bar_label(container)
plt.title('Gender-wise Distribution')
```


```
plt.xticks(rotation=45)
plt.show()
```



In [14]:

```
sns.set_style('whitegrid')
fig, axes = plt.subplots(figsize=(12, 8))
ax = sns.countplot(x='dx', data=meta_data, order = meta_data['dx'].value_counts().index,
palette = 'Paired')
for container in ax.containers:
    ax.bar_label(container)
plt.title('Cell Types Skin Cancer Affected patients')
plt.xticks(rotation=45)
plt.show()
```





In []:

In [15]:

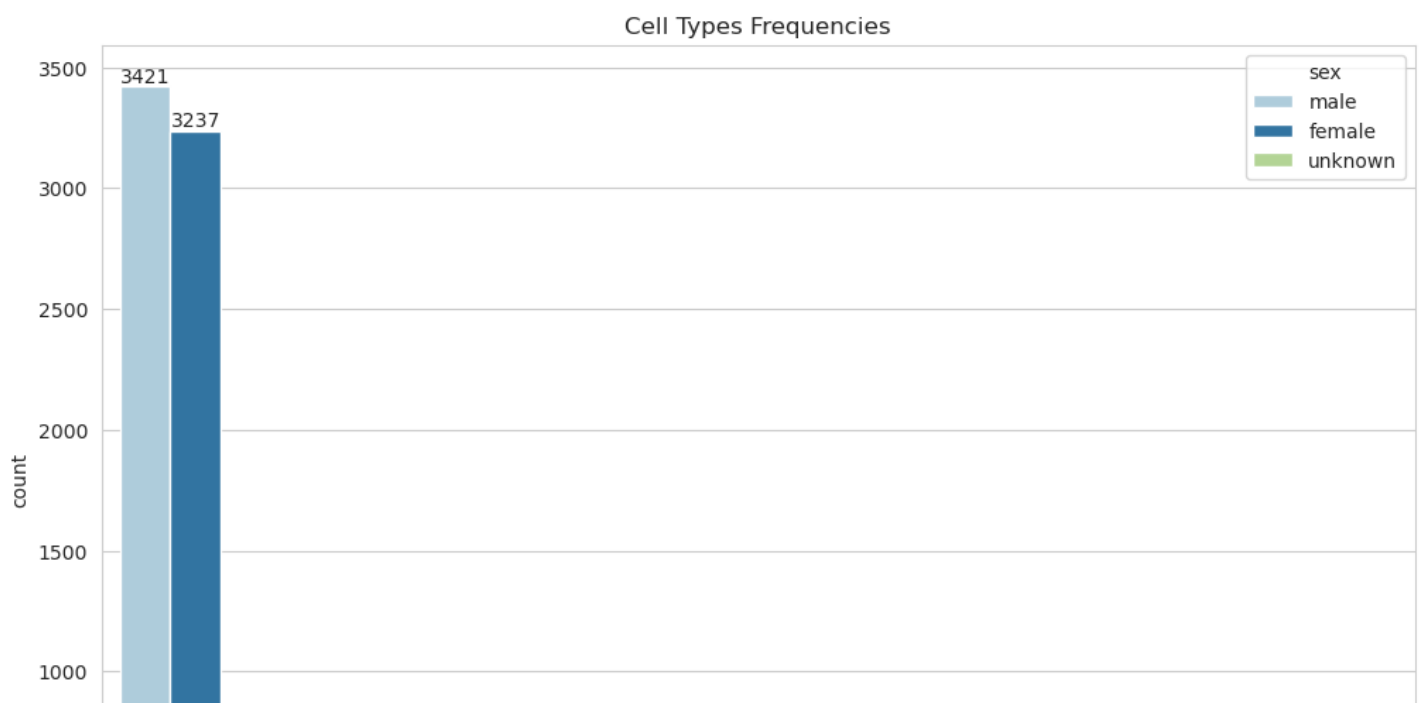
```
classes = {2:'bkl', 4:'nv', 3:'df', 6:'mel', 5:'vasc', 1:'bcc', 0:'akiec'}
```


```
classes_labels=[]
for key in classes.keys():
    classes_labels.append(key)
print(classes_labels)
```

```
[2, 4, 3, 6, 5, 1, 0]
```

In [16]:


```
sns.set_style('whitegrid')
fig, axes = plt.subplots(figsize=(12,8))
ax = sns.countplot(x='dx', hue='sex', data=meta_data, order = meta_data['dx'].value_counts().index, palette = 'Paired')
for container in ax.containers:
    ax.bar_label(container)
plt.title('Cell Types Frequencies')
plt.xticks(rotation=45)
plt.show()
```





In [17]:


```
sns.set_style('whitegrid')
fig, axes = plt.subplots(figsize=(12, 8))
ax = sns.countplot(x='localization', data=meta_data, order = meta_data['localization'].value_counts().index, palette = 'crest')
for container in ax.containers:
    ax.bar_label(container)
plt.title('Localization Area Frequencies')
plt.xticks(rotation=45)
plt.show()
```




In [18]:

```
sns.set_style('whitegrid')
fig, axes = plt.subplots(figsize=(12, 8))
ax = sns.histplot(data=meta_data, x='age')
plt.title('Age Histogram')
plt.show()
```


Age Histogram





In [19]:

```
sns.set_style('whitegrid')
fig, axes = plt.subplots(figsize=(12, 8))
ax = sns.histplot(data=meta_data, x='age', hue='sex', multiple='stack')
plt.title('Age Histogram Gender Wise')
plt.show()
```



In [20]:

```
print(x.shape,y.shape)
# To overcome class imbalance
oversample = RandomOverSampler()
x,y = oversample.fit_resample(x,y)
print(x.shape,y.shape)
```

```
(10015, 2352) (10015,)
(46935, 2352) (46935,)
```

In [21]:

```
# reshaping the data so that it can be taken by convolution neural network(without disturbing the no. of samples)
x = np.array(x).reshape(-1,28,28,3)
print('Shape of X :',x.shape)
print('Shape of y :',y.shape)
```

```
Shape of X : (46935, 28, 28, 3)
Shape of y : (46935,)
```

In [22]:

```
# Splitting Data
X_train, X_test, Y_train, Y_test = train_test_split(x,y, test_size=0.2, random_state=1)
print(X_train.shape,Y_train.shape)
print(X_test.shape , Y_test.shape)
```

```
(37548, 28, 28, 3) (37548,)
(9387, 28, 28, 3) (9387,)
```

In [23]:

```
model_CNN = Sequential()
model_CNN.add(Conv2D(32, kernel_size = (3,3), input_shape = (28, 28, 3), activation = 'relu', padding = 'same'))
model_CNN.add(BatchNormalization())

model_CNN.add(MaxPool2D(pool_size = (2,2)))

model_CNN.add(Conv2D(64, kernel_size = (3,3), activation = 'relu', padding = 'same'))
model_CNN.add(BatchNormalization())

model_CNN.add(MaxPool2D(pool_size = (2,2), padding = 'same'))
model_CNN.add(Conv2D(128, kernel_size = (3,3), activation = 'relu', padding = 'same'))
model_CNN.add(BatchNormalization())

model_CNN.add(MaxPool2D(pool_size = (2,2), padding = 'same'))

model_CNN.add(Flatten())

model_CNN.add(Dense(64, activation = 'relu'))
model_CNN.add(BatchNormalization())

model_CNN.add(Dense(32))
model_CNN.add(Activation(activation='relu'))
model_CNN.add(BatchNormalization())

model_CNN.add(Dense(16))
model_CNN.add(Activation(activation='relu'))
model_CNN.add(BatchNormalization())

model_CNN.add(Dense(7))
model_CNN.add(Activation(activation='softmax'))

optimizer = tf.keras.optimizers.Adam(learning_rate = 0.001)

model_CNN.compile(loss = 'sparse_categorical_crossentropy',
```



```
optimizer = optimizer,
metrics = ['accuracy'])
print(model_CNN.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 32)	896
batch_normalization (Batch Normalization)	(None, 28, 28, 32)	128
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 14, 14, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 14, 14, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_2 (Conv2D)	(None, 7, 7, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 7, 7, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 64)	131136
batch_normalization_3 (Batch Normalization)	(None, 64)	256
dense_1 (Dense)	(None, 32)	2080
activation (Activation)	(None, 32)	0
batch_normalization_4 (Batch Normalization)	(None, 32)	128
dense_2 (Dense)	(None, 16)	528
activation_1 (Activation)	(None, 16)	0
batch_normalization_5 (Batch Normalization)	(None, 16)	64
dense_3 (Dense)	(None, 7)	119
activation_2 (Activation)	(None, 7)	0
Total params: 228,455		
Trainable params: 227,783		
Non-trainable params: 672		
None		

In [24]:

```
from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping
early_stop = EarlyStopping(monitor='val_loss', patience=10, verbose=1, mode='auto')
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=3, verbose=1, mode='auto')
history = model_CNN.fit(X_train,
                        Y_train,
```

```
validation_split=0.2,  
batch_size = 64,  
epochs = 50,  
callbacks = [reduce_lr, early_stop])
```

Epoch 1/50

470/470 [=====] - 15s 10ms/step - loss: 0.9542 - accuracy: 0.6635 - val_loss: 0.9318 - val_accuracy: 0.6377 - lr: 0.0010

Epoch 2/50

470/470 [=====] - 4s 8ms/step - loss: 0.4104 - accuracy: 0.8578 - val_loss: 0.8866 - val_accuracy: 0.6979 - lr: 0.0010

Epoch 3/50

470/470 [=====] - 4s 8ms/step - loss: 0.2900 - accuracy: 0.9001 - val_loss: 0.6961 - val_accuracy: 0.7650 - lr: 0.0010

Epoch 4/50

470/470 [=====] - 4s 8ms/step - loss: 0.1956 - accuracy: 0.9325 - val_loss: 0.4986 - val_accuracy: 0.8198 - lr: 0.0010

Epoch 5/50

470/470 [=====] - 4s 8ms/step - loss: 0.1521 - accuracy: 0.9478 - val_loss: 0.4333 - val_accuracy: 0.8583 - lr: 0.0010

Epoch 6/50

470/470 [=====] - 4s 8ms/step - loss: 0.1286 - accuracy: 0.9552 - val_loss: 0.4032 - val_accuracy: 0.8720 - lr: 0.0010

Epoch 7/50

470/470 [=====] - 4s 8ms/step - loss: 0.0947 - accuracy: 0.9682 - val_loss: 0.1674 - val_accuracy: 0.9463 - lr: 0.0010

Epoch 8/50

470/470 [=====] - 4s 8ms/step - loss: 0.0844 - accuracy: 0.9713 - val_loss: 1.3121 - val_accuracy: 0.6088 - lr: 0.0010

Epoch 9/50

470/470 [=====] - 4s 8ms/step - loss: 0.1012 - accuracy: 0.9646 - val_loss: 0.8057 - val_accuracy: 0.7658 - lr: 0.0010

Epoch 10/50

467/470 [=====>.] - ETA: 0s - loss: 0.0704 - accuracy: 0.9767

Epoch 10: ReduceLROnPlateau reducing learning rate to 0.00010000000474974513.

470/470 [=====] - 4s 9ms/step - loss: 0.0705 - accuracy: 0.9767 - val_loss: 0.2393 - val_accuracy: 0.9178 - lr: 0.0010

Epoch 11/50

470/470 [=====] - 4s 9ms/step - loss: 0.0214 - accuracy: 0.9943 - val_loss: 0.0746 - val_accuracy: 0.9794 - lr: 1.0000e-04

Epoch 12/50

470/470 [=====] - 4s 8ms/step - loss: 0.0119 - accuracy: 0.9976 - val_loss: 0.0778 - val_accuracy: 0.9798 - lr: 1.0000e-04

Epoch 13/50

470/470 [=====] - 4s 8ms/step - loss: 0.0093 - accuracy: 0.9982 - val_loss: 0.0703 - val_accuracy: 0.9811 - lr: 1.0000e-04

Epoch 14/50

470/470 [=====] - 4s 8ms/step - loss: 0.0064 - accuracy: 0.9993 - val_loss: 0.0760 - val_accuracy: 0.9796 - lr: 1.0000e-04

Epoch 15/50

470/470 [=====] - 4s 8ms/step - loss: 0.0048 - accuracy: 0.9996 - val_loss: 0.0708 - val_accuracy: 0.9822 - lr: 1.0000e-04

Epoch 16/50

469/470 [=====>.] - ETA: 0s - loss: 0.0049 - accuracy: 0.9993

Epoch 16: ReduceLROnPlateau reducing learning rate to 1.0000000474974514e-05.

470/470 [=====] - 4s 8ms/step - loss: 0.0049 - accuracy: 0.9993 - val_loss: 0.0831 - val_accuracy: 0.9794 - lr: 1.0000e-04

Epoch 17/50

470/470 [=====] - 4s 9ms/step - loss: 0.0039 - accuracy: 0.9996 - val_loss: 0.0709 - val_accuracy: 0.9816 - lr: 1.0000e-05

Epoch 18/50

470/470 [=====] - 4s 9ms/step - loss: 0.0034 - accuracy: 0.9997 - val_loss: 0.0718 - val_accuracy: 0.9818 - lr: 1.0000e-05

Epoch 19/50

470/470 [=====] - ETA: 0s - loss: 0.0035 - accuracy: 0.9997

Epoch 19: ReduceLROnPlateau reducing learning rate to 1.0000000656873453e-06.

470/470 [=====] - 4s 9ms/step - loss: 0.0035 - accuracy: 0.9997 - val_loss: 0.0708 - val_accuracy: 0.9820 - lr: 1.0000e-05

Epoch 20/50

470/470 [=====] - 4s 8ms/step - loss: 0.0030 - accuracy: 0.9998 - val_loss: 0.0717 - val_accuracy: 0.9816 - lr: 1.0000e-06

Epoch 21/50

```

470/470 [=====] - 4s 8ms/step - loss: 0.0031 - accuracy: 0.9998
- val_loss: 0.0714 - val_accuracy: 0.9812 - lr: 1.0000e-06
Epoch 22/50
469/470 [=====>.] - ETA: 0s - loss: 0.0028 - accuracy: 0.9999
Epoch 22: ReduceLROnPlateau reducing learning rate to 1.0000001111620805e-07.
470/470 [=====] - 4s 8ms/step - loss: 0.0028 - accuracy: 0.9999
- val_loss: 0.0708 - val_accuracy: 0.9819 - lr: 1.0000e-06
Epoch 23/50
470/470 [=====] - 4s 8ms/step - loss: 0.0030 - accuracy: 0.9998
- val_loss: 0.0715 - val_accuracy: 0.9818 - lr: 1.0000e-07
Epoch 23: early stopping

```

In [25]:

```

results = model_CNN.evaluate(X_test , Y_test, verbose=0)

print("CNN Model Test Results")
print("      Test Loss: {:.5f}".format(results[0]))
print("      Test Accuracy: {:.2f}%".format(results[1] * 100))

```

```

CNN Model Test Results
      Test Loss: 0.07844
      Test Accuracy: 97.93%


```

In [26]:

```

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()

```




In [27]:

```

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper right')
plt.show()

```



In [28]:

```
from sklearn.metrics import confusion_matrix , classification_report

y_true_CNN = list(Y_test)
y_pred_CNN = model_CNN.predict(X_test)
y_pred_CNN = list(map(lambda x: np.argmax(x), y_pred_CNN))
print("Predicting First Ten Rows:")
print('Y Actual Values :' , y_true_CNN[0:10])
print('Y Predicted Values :' , y_pred_CNN[0:10])
```

```
294/294 [=====] - 1s 2ms/step
Predicting First Ten Rows:
Y Actual Values : [5, 1, 4, 0, 5, 0, 2, 0, 3, 2]
Y Predicted Values : [5, 1, 4, 0, 5, 0, 2, 0, 3, 2]
```


In [29]:


```
cm_CNN = confusion_matrix(y_true_CNN,y_pred_CNN,labels=classes_labels)
print(confusion_matrix(y_true_CNN,y_pred_CNN,labels=classes_labels))
sns.heatmap(cm_CNN, annot = True, fmt='')
```

```
[[1247    5    1    3    2    4    0]
 [  57 1208    4   92    3    9    1]
 [    0    0 1351    0    0    0    0]
 [   4    9    0 1352    0    0    0]
 [    0    0    0    0 1358    0    0]
 [    0    0    0    0    0 1318    0]
 [    0    0    0    0    0    0 1359]]
```

Out[29]:

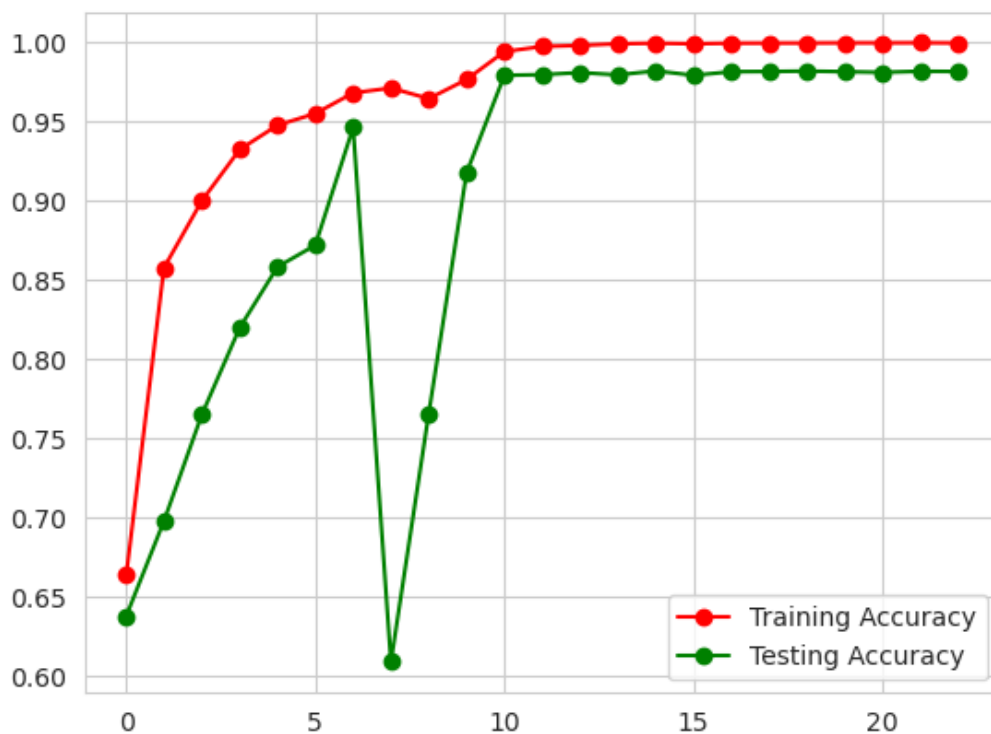
<AxesSubplot:>





In [30]:

```
#training acc vs testing acc graph
plt.plot(history.history["accuracy"], 'ro-', label = "Training Accuracy")
plt.plot(history.history["val_accuracy"], 'go-', label = "Testing Accuracy")
plt.legend()
plt.show()
```



In [31]:

```
#predicting
y_pred_CNN = model_CNN.predict(X_test)
target_names = [f"{classes[i]}" for i in range(7)]
y_pred_CNN = list(map(lambda x: np.argmax(x), y_pred_CNN))
print("CNN Model Prediction Results")
print(classification_report(Y_test, y_pred_CNN, target_names=target_names))
```

294/294 [=====] - 1s 3ms/step

CNN Model Prediction Results

	precision	recall	f1-score	support
akiec	1.00	1.00	1.00	1359
bcc	0.99	1.00	1.00	1318
bkl	0.95	0.99	0.97	1262
df	1.00	1.00	1.00	1351
nv	0.99	0.88	0.93	1374
vasc	1.00	1.00	1.00	1358
mel	0.93	0.99	0.96	1365
accuracy			0.98	9387
macro avg	0.98	0.98	0.98	9387
weighted avg	0.98	0.98	0.98	9387

In [32]:

```
# Layers definitions
from keras import backend as K
for l in range(len(model_CNN.layers)):
    print(l, model_CNN.layers[l])
```

```
0 <keras.layers.convolutional.conv2d.Conv2D object at 0x7fc69fc93390>
1 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7fc69fd3d310>
2 <keras.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7fc6a00f7b50>
3 <keras.layers.convolutional.conv2d.Conv2D object at 0x7fc70e4c4e90>
4 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7fc6a0113110>
5 <keras.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7fc6a00f7950>
6 <keras.layers.convolutional.conv2d.Conv2D object at 0x7fc6a010db10>
7 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7fc6a0105110>
8 <keras.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7fc69ff02390>
9 <keras.layers.resizing.flatten.Flatten object at 0x7fc69fef5250>
10 <keras.layers.core.dense.Dense object at 0x7fc6a0105e90>
11 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7fc6a0105690>
12 <keras.layers.core.dense.Dense object at 0x7fc69ff16650>
13 <keras.layers.core.activation.Activation object at 0x7fc6a04ad450>
14 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7fc69fec0d0>
15 <keras.layers.core.dense.Dense object at 0x7fc69ff0a0d0>
16 <keras.layers.core.activation.Activation object at 0x7fc69fef5f10>
17 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7fc69ff16390>
18 <keras.layers.core.dense.Dense object at 0x7fc69ff0fe10>
19 <keras.layers.core.activation.Activation object at 0x7fc69ff02e50>
```

In [33]:

```
model_CNN.layers[-2]
```

Out[33]:

```
<keras.layers.core.dense.Dense at 0x7fc69ff0fe10>
```

In [34]:

```
import os
os.environ["KERAS_BACKEND"] = "tensorflow"
kerasBKED = os.environ["KERAS_BACKEND"]
print(kerasBKED)
```

```
tensorflow
```

Separating Features Layers from the CNN Model

In [35]:

```
import tensorflow as tf
# feature_extractor = tf.keras.Model(inputs=model_CNN.input,
#                                     outputs=model_CNN.get_layer(-2).output)
# output_layers_model = tf.keras.Model(inputs=model_CNN.input, outputs=model_CNN.output)
# cnn_layer_output = model_CNN.layers[-2].output
# cnn_model_features = tf.keras.Model(inputs=model_CNN.input, outputs=cnn_layer_output)
cnn_model_features = tf.keras.Model(inputs=model_CNN.input, outputs=model_CNN.layers[-3].output)
```

Extracting Features from CNN Model

In [36]:

```
# Extract features from input data using the CNN model
X_train_cnn = cnn_model_features.predict(X_train)
X_test_cnn = cnn_model_features.predict(X_test)
```

```
1174/1174 [=====] - 3s 2ms/step
294/294 [=====] - 1s 2ms/step
```

Integrating CNN with SVM Classifier using Grid Search for Best Parameters

In [37]:

```
import numpy as np
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV

parameters = {'kernel':['rbf'],
              'C':[1, 10, 100, 1000],
              'gamma':[1e-3, 1e-4]}
clf = GridSearchCV(SVC(), parameters)
clf.fit(X_train_cnn, Y_train)
# Evaluate the combined CNN-SVM model on a test dataset
svm_accuracy = clf.score(X_test_cnn, Y_test)
print('SVM Accuracy:', svm_accuracy*100)
y_testSVM = clf.predict(X_test_cnn)
```

SVM Accuracy: 98.60445296686908

In [38]:

```
svm_accuracy = clf.score(X_test_cnn, Y_test)
print('SVM Accuracy:', svm_accuracy*100)
svmclf = clf.best_estimator_
print(svmclf)
svmclf.fit(X_train_cnn, Y_train)
print("Accuracy: {0}".format(accuracy_score(Y_test, y_testSVM)*100))
```

SVM Accuracy: 98.60445296686908
SVC(C=10, gamma=0.001)
Accuracy: 98.60445296686908

In [39]:

```
y_testSVM = svmclf.predict(X_test_cnn)
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

print(classification_report(Y_test, y_testSVM, target_names=target_names))
print("Accuracy: {0}".format(accuracy_score(Y_test, y_testSVM)*100))
```

	precision	recall	f1-score	support
akiec	1.00	1.00	1.00	1359
bcc	0.99	1.00	1.00	1318
bkl	0.97	0.99	0.98	1262
df	1.00	1.00	1.00	1351
nv	0.98	0.93	0.95	1374
vasc	1.00	1.00	1.00	1358
mel	0.96	0.99	0.98	1365
accuracy			0.99	9387
macro avg	0.99	0.99	0.99	9387
weighted avg	0.99	0.99	0.99	9387

Accuracy: 98.60445296686908

In []:

Integrating CNN with Random Forest Classifier using Grid Search for Best Parameters

In [40]:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

parameters = {"max_depth": [3, None],
              "max_features": [1, 3, 10],
              "min_samples_split": [1.0, 3, 10],
              "min_samples_leaf": [1, 3, 10],
              "bootstrap": [True, False],
              "criterion": ["gini", "entropy"],
              "n_estimators": [10, 20, 50]}

rclf = RandomForestClassifier()
rgclf = GridSearchCV(rclf, param_grid=parameters)
rgclf.fit(X_train_cnn, Y_train)
RFC_accuracy = rgclf.score(X_test_cnn, Y_test)
print('Random Forest Classifier Accuracy:', RFC_accuracy*100)
y_test_RF = rgclf.predict(X_test_cnn)
print("Accuracy: {0}".format(accuracy_score(Y_test, y_test_RF)*100))
```

Random Forest Classifier Accuracy: 99.03057419835943
Accuracy: 99.03057419835943

In [41]:

```
y_test_RF = rgclf.predict(X_test_cnn)
print("Accuracy: {0}".format(accuracy_score(Y_test, y_test_RF)*100))
RFclf = rgclf.best_estimator_
print(RFclf)
y_testRFC = RFclf.predict(X_test_cnn)
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

print(classification_report(Y_test, y_testRFC, target_names=target_names))
print("Accuracy: {0}".format(accuracy_score(Y_test, y_testRFC)*100))
```

Accuracy: 99.03057419835943
RandomForestClassifier(max_features=10, min_samples_split=3, n_estimators=20)

	precision	recall	f1-score	support
akiec	1.00	1.00	1.00	1359
bcc	1.00	1.00	1.00	1318
bkl	0.97	1.00	0.99	1262
df	1.00	1.00	1.00	1351
nv	1.00	0.94	0.97	1374
vasc	1.00	1.00	1.00	1358
mel	0.97	0.99	0.98	1365
accuracy			0.99	9387
macro avg	0.99	0.99	0.99	9387
weighted avg	0.99	0.99	0.99	9387

Accuracy: 99.03057419835943

Integrating CNN with KNN Classifier using Grid Search for Best Parameters

In [42]:

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV

parameters = {"n_neighbors": [1, 5, 10, 30],
              "weights": ['uniform', 'distance'],
              "metric": ['minkowski', 'euclidean', 'manhattan'],
```



```

        "algorithm": ['auto', 'ball_tree', 'kd_tree', 'brute']}]
kclf = KNeighborsClassifier()
kgclf = GridSearchCV(kclf, param_grid=parameters)
kgclf.fit(X_train_cnn, Y_train)
KNN_accuracy = kgclf.score(X_test_cnn, Y_test)
print('KNN Classifier Accuracy:', KNN_accuracy*100)

```

KNN Classifier Accuracy: 98.68967721316714

In [43]:

```

y_testKNN = kgclf.predict(X_test_cnn)
KNNclf = kgclf.best_estimator_
print(KNNclf)
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

print(classification_report(Y_test, y_testKNN, target_names=target_names))
print("Accuracy Score: {0}".format(accuracy_score(Y_test, y_testKNN)*100))

```

```

KNeighborsClassifier(n_neighbors=10, weights='distance')

```

	precision	recall	f1-score	support
akiec	1.00	1.00	1.00	1359
bcc	0.99	1.00	1.00	1318
bkl	0.96	1.00	0.98	1262
df	1.00	1.00	1.00	1351
nv	1.00	0.92	0.96	1374
vasc	1.00	1.00	1.00	1358
mel	0.96	1.00	0.98	1365
accuracy			0.99	9387
macro avg	0.99	0.99	0.99	9387
weighted avg	0.99	0.99	0.99	9387

Accuracy Score: 98.68967721316714

Integrating CNN with Logistic Regression Classifier using Grid Search for Best Parameters

In [44]:

```

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

# Create a logistic regression object
lr = LogisticRegression()

# Define the hyperparameter grid to search over
param_grid = {'C': [0.001, 0.01, 0.1, 1, 10, 100],
              'penalty': ['l1', 'l2']}

# Perform grid search with 5-fold cross-validation
grid_search_LR = GridSearchCV(lr, param_grid, cv=5)
grid_search_LR.fit(X_train_cnn, Y_train)

# Print the best hyperparameters and the corresponding accuracy score
print("Best hyperparameters: ", grid_search_LR.best_params_)
y_test_LR = grid_search_LR.predict(X_test_cnn)

print(classification_report(Y_test, y_test_LR, target_names=target_names))
print("Accuracy: {0}".format(accuracy_score(Y_test, y_test_LR)*100))

```

/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence Warning: lbfgs failed to converge (status=1): STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence  
Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:372: FitFailedWarning:  
30 fits failed out of a total of 60.  
The score on these train-test partitions for these parameters will be set to nan.  
If these failures are not expected, you can try to debug them by setting error_score='raise'.  
Below are more details about the failures:
```

Below are more details about the failures:

30 fits failed with the following error:

Traceback (most recent call last):

```
File "/opt/conda/lib/python3.7/site-packages/sklearn/model_selection/_validation.py", line 680, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
File "/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py", line 461, in fit  
    solver = _check_solver(self.solver, self.penalty, self.dual)  
File "/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py", line 449, in _check_solver  
    % (solver, penalty)  
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
/opt/conda/lib/python3.7/site-packages/sklearn/model_selection/_search.py:972: UserWarning: One or more of the test scores are non-finite: [
37768      nan 0.99621788      nan 0.996
      nan 0.99637768
      nan 0.99640432      nan 0.99643095      nan 0.99643095]
category=UserWarning,
```

```
Best hyperparameters: {'C': 10, 'penalty': 'l2'}
      precision    recall  f1-score   support

   akiec         1.00        1.00        1.00        1359
    bcc         0.99        1.00        1.00        1318
    bkl         0.97        0.99        0.98        1262
    df          1.00        1.00        1.00        1351
    nv          0.98        0.92        0.95        1374
   vasc         1.00        1.00        1.00        1358
    mel         0.96        0.99        0.97        1365

 accuracy                   0.99        9387
macro avg         0.99        0.99        0.99        9387
weighted avg         0.99        0.99        0.99        9387
```

Accuracy: 98.51922872057101

```
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence
Warning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
`extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,`

In []:

In []:

In []:

In []: