

Dataset Overview

HAM10000 ("Human Against Machine with 10000 training images") dataset - a large collection of multi-source dermatoscopic images of pigmented lesions

The dermatoscopic images are collected from different populations, acquired and stored by different modalities. The final dataset consists of 10015 dermatoscopic images.

It has 7 different classes of skin cancer which are listed below :

- Melanocytic nevi
- Melanoma
- Benign keratosis-like lesions
- Basal cell carcinoma
- Actinic keratoses
- Vascular lesions
- Dermatofibroma

Importing libraries

In [31]:

```
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
from imblearn.over_sampling import RandomOverSampler
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score, KFold

import os, cv2
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, Flatten, Dense, MaxPool2D, Activation, Batch
Normalization
from sklearn.metrics import classification_report, accuracy_score
```

Reading the Data

In [32]:

```
data = pd.read_csv('/kaggle/input/skin-cancer-mnist-ham10000/hmnist_28_28_RGB.csv')
data.head()
```

Out[32]:


	pixel0000	pixel0001	pixel0002	pixel0003	pixel0004	pixel0005	pixel0006	pixel0007	pixel0008	pixel0009	...	pixel2343	pix
0	192	153	193	195	155	192	197	154	185	202	...	173	
1	25	14	30	68	48	75	123	93	126	158	...	60	
2	192	138	153	200	145	163	201	142	160	206	...	167	
3	38	19	30	95	59	72	143	103	119	171	...	44	
4	158	113	139	194	144	174	215	162	191	225	...	209	

5 rows x 2353 columns



In [33]:

```
f, ax = plt.subplots(2,5)
f.set_size_inches(15, 15)
k = 0
y1 = data['label']
x1 = data.drop(columns = ['label'])
x1= np.array(x1).reshape(-1,28,28,3)
for i in range(2):
    for j in range(5):
        ax[i,j].imshow(x1[k].reshape(28,28,3))
        # ax[i, j].set_title(f"Image {Y_train[4]}")
        k = k + 1
plt.tight_layout()
```



Data Preprocessing

Data Cleaning

In [34]:

```
data['label'].unique()
y = data['label']
x = data.drop(columns = ['label'])
data.isnull().sum().sum() #no null values present
```

Out[34]:

0

In [35]:

```
meta_data = pd.read_csv('/kaggle/input/skin-cancer-mnist-ham10000/HAM10000_metadata.csv')
meta_data.head()
```

Out[35]:

	lesion_id	image_id	dx	dx_type	age	sex	localization
0	HAM_0000118	ISIC_0027419	bkl	histo	80.0	male	scalp
1	HAM_0000118	ISIC_0025030	bkl	histo	80.0	male	scalp
2	HAM_0002730	ISIC_0026769	bkl	histo	80.0	male	scalp
3	HAM_0002730	ISIC_0025661	bkl	histo	80.0	male	scalp
4	HAM_0001466	ISIC_0031633	bkl	histo	75.0	male	ear

In [36]:

```
meta_data['dx'].unique()
```

Out[36]:

```
array(['bkl', 'nv', 'df', 'mel', 'vasc', 'bcc', 'akiec'], dtype=object)
```

In [37]:

```
data.isnull().sum().sum() #no null values present
meta_data.head()
```

Out[37]:

	lesion_id	image_id	dx	dx_type	age	sex	localization
0	HAM_0000118	ISIC_0027419	bkl	histo	80.0	male	scalp
1	HAM_0000118	ISIC_0025030	bkl	histo	80.0	male	scalp
2	HAM_0002730	ISIC_0026769	bkl	histo	80.0	male	scalp
3	HAM_0002730	ISIC_0025661	bkl	histo	80.0	male	scalp
4	HAM_0001466	ISIC_0031633	bkl	histo	75.0	male	ear


Exploratory Data Analysis

In [38]:

```
sns.countplot(x = 'dx', data = meta_data)
plt.xlabel('Disease(Classes)', size=12)
plt.ylabel('Frequency', size=12)
plt.title('Frequency Distribution of Classes')
```

Out[38]:

```
Text(0.5, 1.0, 'Frequency Distribution of Classes')
```



In [39]:


```
sns.set_style('whitegrid')
```

```

colors = ['#87ace8', '#e3784d', 'green']
fig, axes = plt.subplots(figsize=(8, 8))

ax = sns.countplot(x='sex', data=meta_data, palette = 'Paired')
for container in ax.containers:
    ax.bar_label(container)
plt.title('Gender-wise Distribution')
plt.xticks(rotation=45)
plt.show()

```




In [40]:

```

sns.set_style('whitegrid')
fig, axes = plt.subplots(figsize=(12, 8))
ax = sns.countplot(x='dx', data=meta_data, order = meta_data['dx'].value_counts().index,
palette = 'Paired')
for container in ax.containers:
    ax.bar_label(container)
plt.title('Cell Types Skin Cancer Affected patients')
plt.xticks(rotation=45)
plt.show()

```



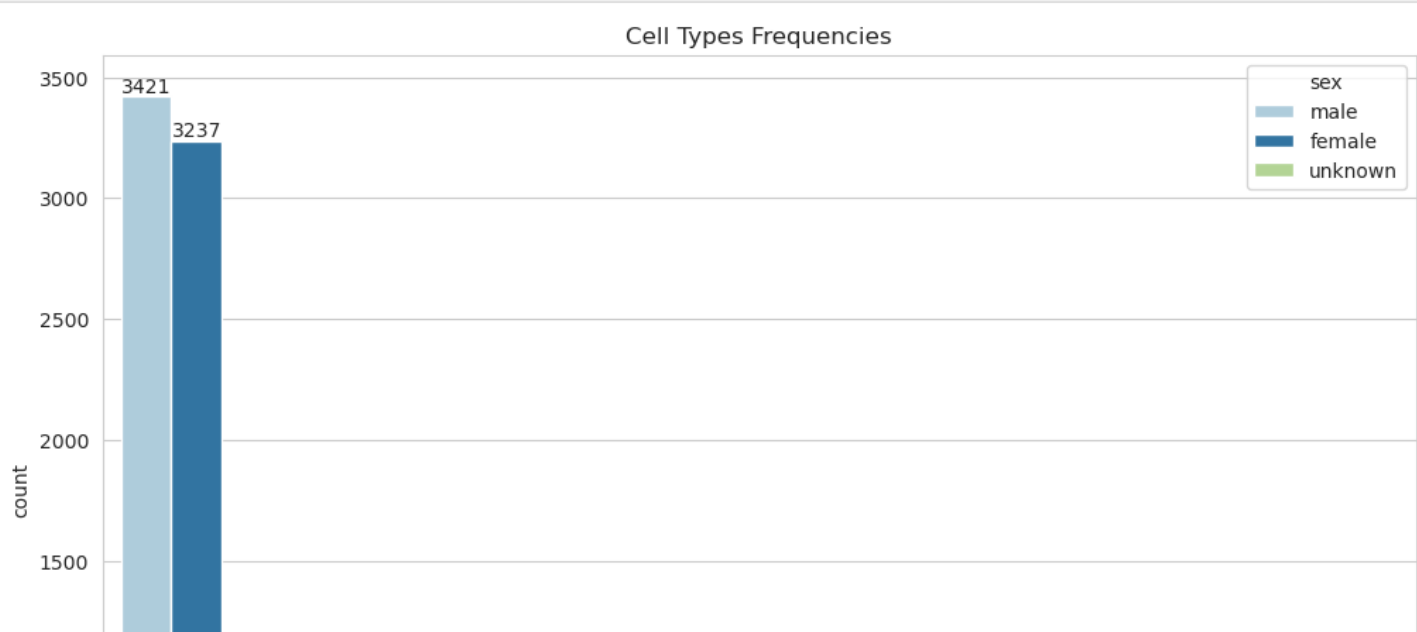



```
In [41]:
classes = {2:'bkl', 4:'nv', 3:'df', 6:'mel', 5:'vasc', 1:'bcc', 0:'akiec'}

classes_labels=[]
for key in classes.keys():
    classes_labels.append(key)
print(classes_labels)

[2, 4, 3, 6, 5, 1, 0]
```


```
In [42]:
sns.set_style('whitegrid')
fig,axes = plt.subplots(figsize=(12,8))
ax = sns.countplot(x='dx',hue='sex', data=meta_data, order = meta_data['dx'].value_counts().index, palette = 'Paired')
for container in ax.containers:
    ax.bar_label(container)
plt.title('Cell Types Frequencies')
plt.xticks(rotation=45)
plt.show()
```





In [43]:


```
sns.set_style('whitegrid')
fig, axes = plt.subplots(figsize=(12, 8))
ax = sns.countplot(x='localization', data=meta_data, order = meta_data['localization'].value_counts().index, palette = 'crest')
for container in ax.containers:
    ax.bar_label(container)
plt.title('Localization Area Frequencies')
plt.xticks(rotation=45)
plt.show()
```




In [44]:

```
sns.set_style('whitegrid')
fig, axes = plt.subplots(figsize=(12, 8))
ax = sns.histplot(data=meta_data, x='age')
plt.title('Age Histogram')
plt.show()
```

Age Histogram



```
In [45]:  
  
sns.set_style('whitegrid')  
fig,axes = plt.subplots(figsize=(12,8))  
ax = sns.histplot(data=meta_data, x='age',hue='sex',multiple='stack')  
plt.title('Age Histogram Gender Wise')  
plt.show()
```



In [46]:

```
print(x.shape,y.shape)
# To overcome class imbalance
oversample = RandomOverSampler()
Data,Label = oversample.fit_resample(x,y)
print(Data.shape,Label.shape)
```

```
(10015, 2352) (10015,)
(46935, 2352) (46935,)
```

In [47]:

```
# reshaping the data so that it can be taken by convolution neural network(without distur
bing the no. of samples)
Data= np.array(Data).reshape(-1,28,28,3)
print('Shape of X :',Data.shape)
print('Shape of y :',Label.shape)
```

```
Shape of X : (46935, 28, 28, 3)
Shape of y : (46935,)
```

In [48]:

```
print(Label)
```

```
0          2
1          2
2          2
3          2
4          2
..
46930      6
46931      6
46932      6
46933      6
46934      6
Name: label, Length: 46935, dtype: int64
```

In [49]:

```
newLabel = []

for i in range(len(Label)):
    for key in classes.keys():
        if Label[i] == key:
            newLabel.append(classes[key])

newLabel[46930]
```

Out[49]:

```
'mel'
```


In [50]:

```
sns.set_style('whitegrid')
fig,axes = plt.subplots(figsize=(8,8))
ax = sns.countplot(x=newLabel,data=data,width=0.5)
for container in ax.containers:
    ax.bar_label(container, fontweight='bold', fontsize=11)
plt.title('Balance Data using Over Sampling', fontweight='bold', fontsize=14)
plt.ylim(0, 8000)
plt.xticks(fontweight='bold', fontsize=12)
plt.yticks(fontweight='bold', fontsize=12)
plt.ylabel("Count",fontweight='bold', fontsize=14)
plt.xlabel("Disease Class", fontweight='bold', fontsize=14)

plt.grid(linewidth=0.8)
plt.show()
```

Balance Data using Over Sampling

Balance Data using Over Sampling



In [51]:

```
Data = np.array(Data).reshape(-1,28,28,3)
print('Shape of X :',Data.shape)
print('Shape of y :',Label.shape)
```

```
Shape of X : (46935, 28, 28, 3)
Shape of y : (46935,)
```

In [52]:

```
# Splitting Data
X_train, X_test, Y_train, Y_test = train_test_split(Data,Label, test_size=0.2, random_state=1)
print(X_train.shape,Y_train.shape)
print(X_test.shape , Y_test.shape)
```

```
(37548, 28, 28, 3) (37548,)
(9387, 28, 28, 3) (9387,)
```

In [53]:

```
classes = {4:'melanocytic nevi',
6: 'melanoma',
2:'benign keratosis-like lesions',
1:'basal cell carcinoma',
5:'pyogenic granulomas and hemorrhage',
0:'Actinic keratoses and intraepithelial carcinomae',
3:'dermatofibroma'}
```

In [54]:

```
Y_train[:10]
```

Out[54]:

```
18033    1
38412    5
26590    2
6910     4
37704    5
15663    0
31556    3
45531    6
24641    2
44902    6
Name: label, dtype: int64
```

In [55]:

```
l_train = np.array(Y_train)
l_train[1]
```

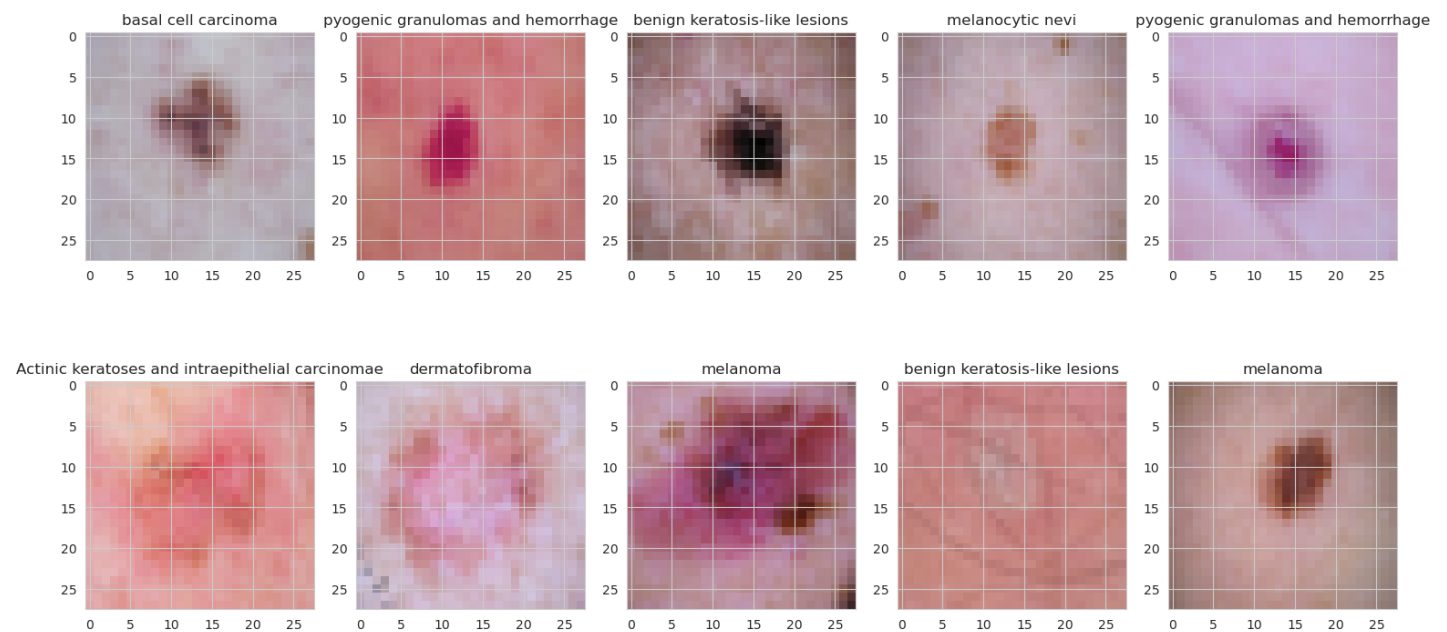
Out[55]:

5

In [56]:

```
f, ax = plt.subplots(2,5)
f.set_size_inches(15, 15)
k = 0
l_train = np.array(Y_train)
for i in range(2):
    for j in range(5):
        plt.subplots_adjust(wspace=5)
        ax[i,j].imshow(X_train[k].reshape(28,28,3))
        ax[i, j].set_title(f"{classes[l_train[k]]}")
        k = k + 1

plt.tight_layout()
```



In [57]:

```
model_CNN = Sequential()
model_CNN.add(Conv2D(32, kernel_size = (3,3), input_shape = (28, 28, 3), activation = 'relu', padding = 'same'))
model_CNN.add(BatchNormalization())
```

```
model_CNN.add(MaxPool2D(pool_size = (2,2)))

model_CNN.add(Conv2D(64, kernel_size = (3,3), activation = 'relu', padding = 'same'))
model_CNN.add(BatchNormalization())

model_CNN.add(MaxPool2D(pool_size = (2,2), padding = 'same'))
model_CNN.add(Conv2D(128, kernel_size = (3,3), activation = 'relu', padding = 'same'))
model_CNN.add(BatchNormalization())

model_CNN.add(MaxPool2D(pool_size = (2,2), padding = 'same'))

model_CNN.add(Flatten())

model_CNN.add(Dense(64, activation = 'relu'))
model_CNN.add(BatchNormalization())

model_CNN.add(Dense(32))
model_CNN.add(Activation(activation='relu'))
model_CNN.add(BatchNormalization())

model_CNN.add(Dense(16))
model_CNN.add(Activation(activation='relu'))
model_CNN.add(BatchNormalization())

model_CNN.add(Dense(7))
model_CNN.add(Activation(activation='softmax'))

optimizer = tf.keras.optimizers.Adam(learning_rate = 0.001)

model_CNN.compile(loss = 'sparse_categorical_crossentropy',
                  optimizer = optimizer,
                  metrics = ['accuracy'])
print(model_CNN.summary())
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 28, 28, 32)	896
batch_normalization_6 (Batch Normalization)	(None, 28, 28, 32)	128
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_4 (Conv2D)	(None, 14, 14, 64)	18496
batch_normalization_7 (Batch Normalization)	(None, 14, 14, 64)	256
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_5 (Conv2D)	(None, 7, 7, 128)	73856
batch_normalization_8 (Batch Normalization)	(None, 7, 7, 128)	512
max_pooling2d_5 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten_1 (Flatten)	(None, 2048)	0
dense_4 (Dense)	(None, 64)	131136
batch_normalization_9 (Batch Normalization)	(None, 64)	256

dense_5 (Dense)	(None, 32)	2080
activation_3 (Activation)	(None, 32)	0
batch_normalization_10 (Batch Normalization)	(None, 32)	128
dense_6 (Dense)	(None, 16)	528
activation_4 (Activation)	(None, 16)	0
batch_normalization_11 (Batch Normalization)	(None, 16)	64
dense_7 (Dense)	(None, 7)	119
activation_5 (Activation)	(None, 7)	0

```

=====
Total params: 228,455
Trainable params: 227,783
Non-trainable params: 672

```

None

In [58]:

```

from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping
early_stop = EarlyStopping(monitor='val_loss', patience=10, verbose=1, mode='auto')
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=3, verbose=1, mode='auto')
history = model_CNN.fit(X_train,
                        Y_train,
                        validation_split=0.2,
                        batch_size = 64,
                        epochs = 50,
                        callbacks = [reduce_lr, early_stop])

```

```

Epoch 1/50
470/470 [=====] - 9s 9ms/step - loss: 1.0044 - accuracy: 0.6432
- val_loss: 0.7764 - val_accuracy: 0.7091 - lr: 0.0010
Epoch 2/50
470/470 [=====] - 4s 8ms/step - loss: 0.4632 - accuracy: 0.8415
- val_loss: 0.8838 - val_accuracy: 0.6871 - lr: 0.0010
Epoch 3/50
470/470 [=====] - 4s 8ms/step - loss: 0.2989 - accuracy: 0.8958
- val_loss: 0.6778 - val_accuracy: 0.7487 - lr: 0.0010
Epoch 4/50
470/470 [=====] - 4s 8ms/step - loss: 0.2196 - accuracy: 0.9224
- val_loss: 1.0653 - val_accuracy: 0.6581 - lr: 0.0010
Epoch 5/50
470/470 [=====] - 4s 8ms/step - loss: 0.1689 - accuracy: 0.9410
- val_loss: 0.3811 - val_accuracy: 0.8569 - lr: 0.0010
Epoch 6/50
470/470 [=====] - 5s 10ms/step - loss: 0.1274 - accuracy: 0.9568
- val_loss: 0.5980 - val_accuracy: 0.8057 - lr: 0.0010
Epoch 7/50
470/470 [=====] - 4s 8ms/step - loss: 0.1164 - accuracy: 0.9581
- val_loss: 0.4327 - val_accuracy: 0.8523 - lr: 0.0010
Epoch 8/50
470/470 [=====] - 4s 9ms/step - loss: 0.0936 - accuracy: 0.9674
- val_loss: 0.2524 - val_accuracy: 0.9142 - lr: 0.0010
Epoch 9/50
470/470 [=====] - 4s 9ms/step - loss: 0.0931 - accuracy: 0.9677
- val_loss: 0.5019 - val_accuracy: 0.8316 - lr: 0.0010
Epoch 10/50
470/470 [=====] - 4s 9ms/step - loss: 0.0572 - accuracy: 0.9804
- val_loss: 0.2258 - val_accuracy: 0.9308 - lr: 0.0010
Epoch 11/50
470/470 [=====] - 4s 8ms/step - loss: 0.0639 - accuracy: 0.9777
- val_loss: 0.6433 - val_accuracy: 0.8142 - lr: 0.0010
Epoch 12/50

```

```
470/470 [=====] - 4s 8ms/step - loss: 0.0524 - accuracy: 0.9824
- val_loss: 1.0235 - val_accuracy: 0.7035 - lr: 0.0010
Epoch 13/50
467/470 [=====>.] - ETA: 0s - loss: 0.0779 - accuracy: 0.9731
Epoch 13: ReduceLROnPlateau reducing learning rate to 0.00010000000474974513.
470/470 [=====] - 4s 9ms/step - loss: 0.0779 - accuracy: 0.9731
- val_loss: 0.2557 - val_accuracy: 0.9185 - lr: 0.0010
Epoch 14/50
470/470 [=====] - 4s 9ms/step - loss: 0.0221 - accuracy: 0.9936
- val_loss: 0.0764 - val_accuracy: 0.9776 - lr: 1.0000e-04
Epoch 15/50
470/470 [=====] - 4s 8ms/step - loss: 0.0119 - accuracy: 0.9972
- val_loss: 0.0770 - val_accuracy: 0.9779 - lr: 1.0000e-04
Epoch 16/50
470/470 [=====] - 4s 8ms/step - loss: 0.0082 - accuracy: 0.9984
- val_loss: 0.0731 - val_accuracy: 0.9792 - lr: 1.0000e-04
Epoch 17/50
470/470 [=====] - 4s 9ms/step - loss: 0.0062 - accuracy: 0.9989
- val_loss: 0.0726 - val_accuracy: 0.9799 - lr: 1.0000e-04
Epoch 18/50
470/470 [=====] - 4s 8ms/step - loss: 0.0045 - accuracy: 0.9996
- val_loss: 0.0774 - val_accuracy: 0.9804 - lr: 1.0000e-04
Epoch 19/50
470/470 [=====] - 4s 8ms/step - loss: 0.0038 - accuracy: 0.9995
- val_loss: 0.0765 - val_accuracy: 0.9794 - lr: 1.0000e-04
Epoch 20/50
463/470 [=====>.] - ETA: 0s - loss: 0.0037 - accuracy: 0.9995
Epoch 20: ReduceLROnPlateau reducing learning rate to 1.0000000474974514e-05.
470/470 [=====] - 4s 9ms/step - loss: 0.0040 - accuracy: 0.9994
- val_loss: 0.0843 - val_accuracy: 0.9784 - lr: 1.0000e-04
Epoch 21/50
470/470 [=====] - 4s 9ms/step - loss: 0.0028 - accuracy: 0.9997
- val_loss: 0.0740 - val_accuracy: 0.9808 - lr: 1.0000e-05
Epoch 22/50
470/470 [=====] - 4s 8ms/step - loss: 0.0028 - accuracy: 0.9997
- val_loss: 0.0742 - val_accuracy: 0.9806 - lr: 1.0000e-05
Epoch 23/50
470/470 [=====] - 4s 9ms/step - loss: 0.0028 - accuracy: 0.9997
- val_loss: 0.0723 - val_accuracy: 0.9811 - lr: 1.0000e-05
Epoch 24/50
470/470 [=====] - 4s 10ms/step - loss: 0.0023 - accuracy: 0.9999
- val_loss: 0.0749 - val_accuracy: 0.9794 - lr: 1.0000e-05
Epoch 25/50
470/470 [=====] - 4s 10ms/step - loss: 0.0021 - accuracy: 0.9999
- val_loss: 0.0756 - val_accuracy: 0.9799 - lr: 1.0000e-05
Epoch 26/50
465/470 [=====>.] - ETA: 0s - loss: 0.0021 - accuracy: 0.9999
Epoch 26: ReduceLROnPlateau reducing learning rate to 1.0000000656873453e-06.
470/470 [=====] - 4s 8ms/step - loss: 0.0022 - accuracy: 0.9999
- val_loss: 0.0763 - val_accuracy: 0.9795 - lr: 1.0000e-05
Epoch 27/50
470/470 [=====] - 4s 8ms/step - loss: 0.0020 - accuracy: 0.9999
- val_loss: 0.0772 - val_accuracy: 0.9792 - lr: 1.0000e-06
Epoch 28/50
470/470 [=====] - 4s 8ms/step - loss: 0.0021 - accuracy: 0.9999
- val_loss: 0.0772 - val_accuracy: 0.9792 - lr: 1.0000e-06
Epoch 29/50
469/470 [=====>.] - ETA: 0s - loss: 0.0020 - accuracy: 0.9998
Epoch 29: ReduceLROnPlateau reducing learning rate to 1.0000001111620805e-07.
470/470 [=====] - 4s 9ms/step - loss: 0.0020 - accuracy: 0.9998
- val_loss: 0.0757 - val_accuracy: 0.9796 - lr: 1.0000e-06
Epoch 30/50
470/470 [=====] - 4s 8ms/step - loss: 0.0022 - accuracy: 0.9997
- val_loss: 0.0758 - val_accuracy: 0.9798 - lr: 1.0000e-07
Epoch 31/50
470/470 [=====] - 4s 8ms/step - loss: 0.0021 - accuracy: 0.9999
- val_loss: 0.0760 - val_accuracy: 0.9798 - lr: 1.0000e-07
Epoch 32/50
466/470 [=====>.] - ETA: 0s - loss: 0.0020 - accuracy: 0.9998
Epoch 32: ReduceLROnPlateau reducing learning rate to 1.000000082740371e-08.
470/470 [=====] - 4s 8ms/step - loss: 0.0020 - accuracy: 0.9998
- val_loss: 0.0770 - val_accuracy: 0.9796 - lr: 1.0000e-07
```

Epoch 33/50
470/470 [=====] - 4s 8ms/step - loss: 0.0021 - accuracy: 0.9999
- val_loss: 0.0758 - val_accuracy: 0.9800 - lr: 1.0000e-08
Epoch 33: early stopping

In [59]:


```
results = model_CNN.evaluate(X_test , Y_test, verbose=0)

print("CNN Model Test Results")
print("      Test Loss: {:.5f}".format(results[0]))
print("      Test Accuracy: {:.2f}%".format(results[1] * 100))
```

CNN Model Test Results
 Test Loss: 0.07440
 Test Accuracy: 98.08%

In [60]:


```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```



In [61]:

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper right')
plt.show()
```





In [62]:

```
from sklearn.metrics import confusion_matrix , classification_report

y_true_CNN = list(Y_test)
y_pred_CNN = model_CNN.predict(X_test)
y_pred_CNN = list(map(lambda x: np.argmax(x), y_pred_CNN))
print("Predicting First Ten Rows:")
print('Y Actual Values :' , y_true_CNN[0:10])
print('Y Predicted Values :' , y_pred_CNN[0:10])
```

```
294/294 [=====] - 1s 2ms/step
Predicting First Ten Rows:
Y Actual Values : [5, 1, 4, 0, 5, 0, 2, 0, 3, 2]
Y Predicted Values : [5, 1, 4, 0, 5, 0, 2, 0, 3, 2]
```

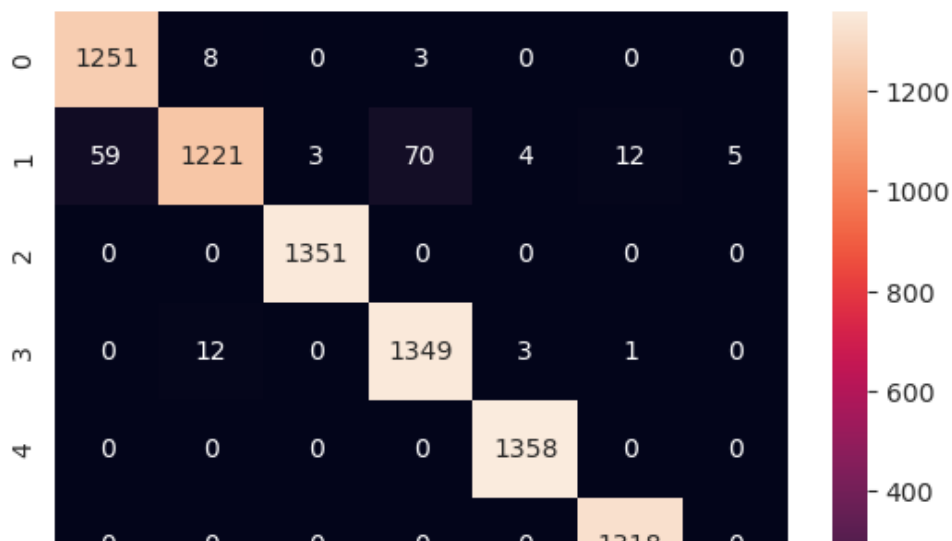
In [63]:

```
cm_CNN = confusion_matrix(y_true_CNN,y_pred_CNN,labels=classes_labels)
print(confusion_matrix(y_true_CNN,y_pred_CNN,labels=classes_labels))
sns.heatmap(cm_CNN, annot = True, fmt='')
```

```
[[1251    8    0    3    0    0    0]
 [  59 1221    3   70    4   12    5]
 [   0    0 1351    0    0    0    0]
 [   0   12    0 1349    3    1    0]
 [   0    0    0    0 1358    0    0]
 [   0    0    0    0    0 1318    0]
 [   0    0    0    0    0    0 1359]]
```

Out[63]:


<AxesSubplot:>





In [64]:

```
#training acc vs testing acc graph
plt.plot(history.history["accuracy"] , 'ro-' , label = "Training Accuracy")
plt.plot(history.history["val_accuracy"] , 'go-' , label = "Testing Accuracy")
plt.legend()
plt.show()
```



In [65]:

```
#predicting
y_pred_CNN = model_CNN.predict(X_test)
target_names = [f"{classes[i]}" for i in range(7)]
y_pred_CNN = list(map(lambda x: np.argmax(x), y_pred_CNN))
print("CNN Model Prediction Results")
print(classification_report(Y_test , y_pred_CNN,target_names=target_names))
```

294/294 [=====] - 1s 2ms/step

CNN Model Prediction Results

	precision	recall	f1-score	support
t				
Actinic keratoses and intraepithelial carcinomae	1.00	1.00	1.00	135
basal cell carcinoma	0.99	1.00	1.00	131
benign keratosis-like lesions	0.95	0.99	0.97	126
dermatofibroma	1.00	1.00	1.00	135
melanocytic nevi	0.98	0.89	0.93	137
pyogenic granulomas and hemorrhage	0.99	1.00	1.00	135
melanoma	0.95	0.99	0.97	136
accuracy			0.98	93

7	macro avg	0.98	0.98	0.98	938
7	weighted avg	0.98	0.98	0.98	938

In [66]:

```
# Layers definitions
from keras import backend as K
for l in range(len(model_CNN.layers)):
    print(l, model_CNN.layers[l])

0 <keras.layers.convolutional.conv2d.Conv2D object at 0x7be998003350>
1 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7be938657e90>
2 <keras.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7be997db88d0>
3 <keras.layers.convolutional.conv2d.Conv2D object at 0x7be938637310>
4 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7be938623810>
5 <keras.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7be9385cd690>
6 <keras.layers.convolutional.conv2d.Conv2D object at 0x7be9385c6cd0>
7 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7be9385d7e50>
8 <keras.layers.pooling.max_pooling2d.MaxPooling2D object at 0x7be9981c45d0>
9 <keras.layers.resizing.flatten.Flatten object at 0x7be9385cd3d0>
10 <keras.layers.core.dense.Dense object at 0x7be9385fb310>
11 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7be9385e71d0>
12 <keras.layers.core.dense.Dense object at 0x7be9385987d0>
13 <keras.layers.core.activation.Activation object at 0x7be9385f5f50>
14 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7be9385fb510>
15 <keras.layers.core.dense.Dense object at 0x7be938581090>
16 <keras.layers.core.activation.Activation object at 0x7be9385b7a90>
17 <keras.layers.normalization.batch_normalization.BatchNormalization object at 0x7be938598790>
18 <keras.layers.core.dense.Dense object at 0x7be938581050>
19 <keras.layers.core.activation.Activation object at 0x7be938560e90>
```

In [67]:

```
model_CNN.layers[-2]
```

Out[67]:

```
<keras.layers.core.dense.Dense at 0x7be938581050>
```

In [68]:

```
import os
os.environ["KERAS_BACKEND"] = "tensorflow"
kerasBKED = os.environ["KERAS_BACKEND"]
print(kerasBKED)
```

```
tensorflow
```

Separating Features Layers from the CNN Model

In [69]:

```
import tensorflow as tf
# feature_extractor = tf.keras.Model(inputs=model_CNN.input,
#                                     outputs=model_CNN.get_layer(-2).output)
# output_layers_model = tf.keras.Model(inputs=model_CNN.input, outputs=model_CNN.output)
# cnn_layer_output = model_CNN.layers[-2].output
# cnn_model_features = tf.keras.Model(inputs=model_CNN.input, outputs=cnn_layer_output)
cnn_model_features = tf.keras.Model(inputs=model_CNN.input, outputs=model_CNN.layers[-3].output)
```

Extracting Features from CNN Model

In [70]:

```
# Extract features from input data using the CNN model
X_train_cnn = cnn_model_features.predict(X_train)
X_test_cnn = cnn_model_features.predict(X_test)
```

1174/1174 [=====] - 3s 2ms/step
294/294 [=====] - 1s 2ms/step

Integrating CNN with SVM Classifier using Grid Search for Best Parameters

In [71]:

```
import numpy as np
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV

clf=SVC(kernel='rbf',C=100,gamma=0.001)
kf=KFold(n_splits=5)
scores_SVM = cross_val_score(clf, X_train_cnn, Y_train, cv=kf)
print(scores_SVM)
print("%0.2f accuracy with a standard deviation of %0.2f" % (scores_SVM.mean(), scores_SVM.std()))
clf.fit(X_train_cnn, Y_train)
# Evaluate the combined CNN-SVM model on a test dataset
svm_accuracy = clf.score(X_test_cnn, Y_test)
print('SVM Accuracy:', svm_accuracy*100)
y_testSVM = clf.predict(X_test_cnn)
```

[1. 0.99986684 1. 1. 0.98122253]
1.00 accuracy with a standard deviation of 0.01
SVM Accuracy: 98.87077873655055

In [72]:

```
svm_accuracy = clf.score(X_test_cnn, Y_test)
print('SVM Accuracy:', svm_accuracy*100)
```

SVM Accuracy: 98.87077873655055

In [73]:

```
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

print(classification_report(Y_test, y_testSVM,target_names=target_names))
print("Accuracy: {0}".format(accuracy_score(Y_test, y_testSVM)*100))
```

	precision	recall	f1-score	suppor
t				
Actinic keratoses and intraepithelial carcinomae	1.00	1.00	1.00	135
9				
basal cell carcinoma	1.00	1.00	1.00	131
8				
benign keratosis-like lesions	0.98	0.99	0.98	126
2				
dermatofibroma	1.00	1.00	1.00	135
1				
melanocytic nevi	0.98	0.94	0.96	137
4				
pyogenic granulomas and hemorrhage	1.00	1.00	1.00	135
8				
melanoma	0.97	0.99	0.98	136
5				
accuracy			0.99	93

87					
	macro avg	0.99	0.99	0.99	938
7					
	weighted avg	0.99	0.99	0.99	938
7					

Accuracy: 98.87077873655055

Integrating CNN with Random Forest Classifier using Grid Search for Best Parameters

In [74]:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
rgclf = RandomForestClassifier(max_depth=3,
                              max_features=1,
                              min_samples_split=3,
                              bootstrap=False,
                              criterion="entropy",
                              n_estimators=20)
kf=KFold(n_splits=5)
scores_RF = cross_val_score(rgclf, X_train_cnn, Y_train, cv=kf)
print(scores_RF)
print("%0.2f accuracy with a standard deviation of %0.2f" % (scores_RF.mean(), scores_RF.
std()))
rgclf.fit(X_train_cnn, Y_train)
RFC_accuracy = rgclf.score(X_test_cnn, Y_test)
print('Random Forest Classifier Accuracy:', RFC_accuracy*100)
y_test_RF = rgclf.predict(X_test_cnn)
print("Accuracy: {0}".format(accuracy_score(Y_test, y_test_RF)*100))
```

[0.99720373 0.9976032 0.99533955 0.99826874 0.96990278]
0.99 accuracy with a standard deviation of 0.01
Random Forest Classifier Accuracy: 96.65494833280069
Accuracy: 96.65494833280069

In [75]:

```
y_test_RF = rgclf.predict(X_test_cnn)
print("Accuracy: {0}".format(accuracy_score(Y_test, y_test_RF)*100))
```

Accuracy: 96.65494833280069

In [76]:

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV

# parameters = {"n_neighbors": [1],
#               "weights": ['uniform', 'distance'],
#               "metric": ['minkowski', 'euclidean', 'manhattan'],
#               "algorithm": ['auto', 'ball_tree', 'kd_tree', 'brute']}
kgclf = KNeighborsClassifier(n_neighbors=1,
                             weights='distance',
                             metric='euclidean',
                             algorithm='auto')
kf=KFold(n_splits=5)
scores_KNN = cross_val_score(kgclf, X_train_cnn, Y_train, cv=kf)
print("%0.2f accuracy with a standard deviation of %0.2f" % (scores_KNN.mean(), scores_KN
N.std()))
kgclf.fit(X_train_cnn, Y_train)
KNN_accuracy = kgclf.score(X_test_cnn, Y_test)
print('KNN Classifier Accuracy:', KNN_accuracy*100)
```

1.00 accuracy with a standard deviation of 0.01
KNN Classifier Accuracy: 99.12645147544477

In [77]:

```
y_testKNN = kgclf.predict(X_test_cnn)
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

print(classification_report(Y_test, y_testKNN, target_names=target_names))
print("Accuracy Score: {0}".format(accuracy_score(Y_test, y_testKNN)*100))
```

	precision	recall	f1-score	support
t				
Actinic keratoses and intraepithelial carcinomae	1.00	1.00	1.00	135
9				
basal cell carcinoma	0.99	1.00	1.00	131
8				
benign keratosis-like lesions	0.98	1.00	0.99	126
2				
dermatofibroma	1.00	1.00	1.00	135
1				
melanocytic nevi	1.00	0.94	0.97	137
4				
pyogenic granulomas and hemorrhage	1.00	1.00	1.00	135
8				
melanoma	0.97	1.00	0.99	136
5				
accuracy			0.99	93
87				
macro avg	0.99	0.99	0.99	938
7				
weighted avg	0.99	0.99	0.99	938
7				

Accuracy Score: 99.12645147544477

Integrating CNN with Logistic Regression Classifier using Grid Search for Best Parameters

In [78]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

# Create a logistic regression object
grid_search_LR = LogisticRegression(C=100,
                                    penalty='l2')
kf=KFold(n_splits=5)
scores_LR = cross_val_score(grid_search_LR, X_train_cnn, Y_train, cv=kf)
print(scores_LR)
print("%0.2f accuracy with a standard deviation of %0.2f" % (scores_LR.mean(), scores_LR.
std()))
# Perform grid search with 5-fold cross-validation
grid_search_LR.fit(X_train_cnn, Y_train)

# Print the best hyperparameters and the corresponding accuracy score
y_test_LR = grid_search_LR.predict(X_test_cnn)

print(classification_report(Y_test, y_test_LR, target_names=target_names))
print("Accuracy: {0}".format(accuracy_score(Y_test, y_test_LR)*100))
```

/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence
Warning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase thenumber of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence
Warning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence
Warning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence
Warning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,

[1.	1.	1.	1.	0.98108936]				
1.00	accuracy with a standard deviation of 0.01							
			precision	recall	f1-score	suppor		
t								
Actinic keratoses and intraepithelial carcinomae			1.00	1.00	1.00	135		
9								
		basal cell carcinoma	1.00	1.00	1.00	131		
8								
		benign keratosis-like lesions	0.97	0.99	0.98	126		
2								
		dermatofibroma	1.00	1.00	1.00	135		
1								
		melanocytic nevi	0.98	0.94	0.96	137		
4								
		pyogenic granulomas and hemorrhage	1.00	1.00	1.00	135		
8								
		melanoma	0.97	0.99	0.98	136		
5								
		accuracy			0.99	93		
87								
		macro avg	0.99	0.99	0.99	938		
7								
		weighted avg	0.99	0.99	0.99	938		
7								

Accuracy: 98.78555449025248

/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818: Convergence
Warning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,

In []:

In []:

In []: